

Neural Network-Based European Option Pricing under Black-Scholes SDE Dynamics

Project group

May 7, 2026

Abstract

This project studies whether a feed-forward neural network can accurately approximate the Black-Scholes pricing function for European call options. The problem is addressed in a controlled setting: data are generated synthetically from the analytical Black-Scholes formula, while Monte Carlo simulation is used as a numerical benchmark. The neural network is trained in a supervised learning framework to learn the mapping from market and contract parameters to the theoretical option price. The goal is not to build a real trading system, but to evaluate the role of a neural network as a pricing surrogate in a mathematically well-defined financial setting.

Contents

Abstract	ii
1 Introduction	1
2 Theoretical Framework	2
2.1 Underlying Asset Dynamics	2
2.2 European Call Option	2
2.3 Black-Scholes Formula	2
2.4 Monte Carlo Pricing	2
3 Methodology	4
3.1 Synthetic Dataset	4
3.2 Neural Network	4
3.3 Training and Validation	4
3.4 Metrics	5
4 Implementation	6
4.1 Code Structure	6
4.2 Computational Efficiency	6
4.3 Reproducibility	6
5 Experiments	7
5.1 Final Experiment Design	7
5.2 Reproducible Artifacts	7
5.3 Evaluation Protocol	7
5.4 Diagnostic Visualizations	8
6 Results	9
6.1 Neural Network Approximation	9
6.2 Error Diagnostics	10
6.3 Monte Carlo Benchmark	11
6.4 Discussion	11
7 Conclusions and Next Steps	12

1 Introduction

Derivative pricing is one of the classical problems in quantitative finance. For European options, the Black-Scholes model provides a closed-form price under precise assumptions on the dynamics of the underlying asset, the market, and volatility [1]. The availability of an analytical solution makes Black-Scholes an ideal controlled environment for evaluating the accuracy of numerical methods and machine learning models.

The research question of the project is:

Can a feed-forward neural network accurately approximate the Black-Scholes pricing function for European call options?

The project compares three approaches:

- the analytical Black-Scholes formula, used as the ground truth;
- Monte Carlo simulation, used as a numerical benchmark;
- a feed-forward neural network, used as a supervised learning model and pricing surrogate.

The main objective is not to replace Black-Scholes in a setting where a closed-form formula is available, but to verify whether a neural network can learn a known pricing function. This is a useful first step for understanding neural pricing methods in more complex settings, where analytical formulas are unavailable or computationally expensive.

2 Theoretical Framework

2.1 Underlying Asset Dynamics

In the Black-Scholes model, the price of the underlying asset follows a geometric Brownian motion under the risk-neutral measure. This formulation is standard in continuous-time financial market modeling [4]:

$$dS_t = rS_t dt + \sigma S_t dW_t, \quad (2.1)$$

where S_t is the asset price at time t , r is the risk-free rate, σ is the constant volatility, and W_t is a standard Brownian motion.

The solution of the SDE is:

$$S_T = S_0 \exp \left[\left(r - \frac{1}{2} \sigma^2 \right) T + \sigma \sqrt{T} Z \right], \quad Z \sim \mathcal{N}(0, 1). \quad (2.2)$$

2.2 European Call Option

A European call option with strike K and maturity T has payoff:

$$\max(S_T - K, 0). \quad (2.3)$$

Under risk-neutral pricing, its value is the discounted expected payoff:

$$C = e^{-rT} \mathbb{E}^{\mathbb{Q}} [\max(S_T - K, 0)]. \quad (2.4)$$

2.3 Black-Scholes Formula

The analytical price of a European call option is:

$$C_{BS} = S_0 N(d_1) - K e^{-rT} N(d_2), \quad (2.5)$$

where $N(\cdot)$ is the cumulative distribution function of a standard normal random variable and:

$$d_1 = \frac{\ln(S_0/K) + (r + \frac{1}{2}\sigma^2) T}{\sigma \sqrt{T}}, \quad (2.6)$$

$$d_2 = d_1 - \sigma \sqrt{T}. \quad (2.7)$$

In this project, C_{BS} is used both as the supervised learning target and as the reference value for evaluation.

2.4 Monte Carlo Pricing

Monte Carlo simulation approximates the risk-neutral expectation by generating many realizations of S_T :

$$\hat{C}_{MC} = e^{-rT} \frac{1}{M} \sum_{i=1}^M \max(S_T^{(i)} - K, 0). \quad (2.8)$$

For sufficiently large M , $\hat{C}_{MC}$ converges to C_{BS} . In this project, Monte Carlo is used as a numerical benchmark and as a comparison point for the neural network.

3 Methodology

3.1 Synthetic Dataset

The dataset is generated by randomly sampling option and market parameters. Each observation contains:

$$x = (S_0, K, T, r, \sigma), \quad (3.1)$$

and the target is:

$$y = C_{BS}(S_0, K, T, r, \sigma). \quad (3.2)$$

The initial parameter ranges are:

Variable	Description	Range
S_0	initial underlying price	[50, 150]
K	strike price	[50, 150]
T	maturity in years	[0.1, 2]
r	risk-free rate	[0, 0.05]
σ	volatility	[0.1, 0.6]

Table 3.1: Parameter ranges used to generate the synthetic dataset.

Using synthetic data makes the problem controlled and allows the neural network error to be evaluated directly against the exact target function.

3.2 Neural Network

The model is a feed-forward neural network for regression. This choice is motivated by the fact that sufficiently large feed-forward networks are universal function approximators under general assumptions [2, 3]. The initial architecture is:

$$5 \rightarrow 64 \rightarrow 64 \rightarrow 32 \rightarrow 1, \quad (3.3)$$

with ReLU activation functions in the hidden layers. The network approximates the function:

$$f_{\theta}(S_0, K, T, r, \sigma) \approx C_{BS}. \quad (3.4)$$

The optimized loss is the mean squared error:

$$\mathcal{L}(\theta) = \frac{1}{n} \sum_{i=1}^n (f_{\theta}(x_i) - y_i)^2. \quad (3.5)$$

3.3 Training and Validation

The data are split into training, validation, and test sets. Inputs are standardized using **StandardScaler**; the target is also scaled during training to improve numerical stability and convergence. The network is trained with Adam, early stopping, and a fixed random seed for reproducibility.

3.4 Metrics

The main metrics are:

- Mean Absolute Error (MAE);
- Root Mean Squared Error (RMSE);
- coefficient of determination R^2 ;
- MAPE computed only on prices greater than 1, to avoid instability when the theoretical option price is close to zero.

4 Implementation

The project is implemented in Python with a modular structure. The main libraries are NumPy, pandas, SciPy, scikit-learn, matplotlib, and PyTorch.

4.1 Code Structure

Listing 4.1: Main project structure.

```
src/nn_option_pricing/  
  black_scholes.py  # analytical formula and payoff  
  monte_carlo.py   # efficient Monte Carlo simulation  
  dataset.py       # synthetic data generation and storage  
  model.py         # PyTorch feed-forward neural network  
  training.py      # training loop, scaling, early stopping  
  evaluation.py    # regression metrics  
  plots.py        # visualizations  
  pipeline.py      # end-to-end experiment  
scripts/  
  run_experiment.py # command-line entry point
```

4.2 Computational Efficiency

The Black-Scholes formula and the dataset generation procedure are implemented in vectorized form with NumPy. The Monte Carlo method uses the exact terminal distribution of the geometric Brownian motion at time T , avoiding unnecessary time discretization. Moreover, Monte Carlo simulation is batched both over options and over simulated paths, which keeps memory usage under control.

The neural network training uses PyTorch **DataLoader**, mini-batch gradient descent, input standardization, and target scaling. These choices make the code suitable both for quick development runs and for larger experiments.

4.3 Reproducibility

To make the experiments reproducible, the project fixes the random seeds of NumPy, Python, and PyTorch. The artifacts produced by an experiment include:

- the generated synthetic dataset;
- the model checkpoint;
- the input scaler;
- the target scaler;
- metrics in JSON format;
- diagnostic plots.

5 Experiments

5.1 Final Experiment Design

The final experiment was designed to evaluate the neural network as a supervised surrogate of the Black-Scholes European call pricing function on a sufficiently large synthetic dataset. The analytical Black-Scholes formula is used as the reference target, while Monte Carlo simulation is used as a numerical benchmark on a deterministic subset of the test set.

The final run used 100,000 synthetic option contracts sampled from the parameter ranges defined in the methodology. The neural network was trained for up to 200 epochs with early stopping, using mini-batches of size 1024. The Monte Carlo benchmark used 50,000 simulated paths on 512 test options.

Listing 5.1: Command used for the final experiment.

```
.venv/bin/python scripts/run_experiment.py \  
--n-samples 100000 \  
--max-epochs 200 \  
--batch-size 1024 \  
--mc-n-paths 50000 \  
--mc-evaluation-samples 512 \  
--data-dir data/final \  
--output-dir outputs/final
```

5.2 Reproducible Artifacts

The complete generated dataset, trained model checkpoint, and fitted scalars are stored under **data/final/** and **outputs/final/**. These directories are ignored by Git because they can be reproduced from the code and experiment configuration.

The selected final artifacts used in the report are tracked in **results/final/**. This directory contains:

- the experiment configuration snapshot;
- neural network metrics;
- Monte Carlo benchmark metrics;
- diagnostic figures;
- a short verification note reporting the final test and documentation checks.

5.3 Evaluation Protocol

The neural network is evaluated on a held-out test set that is not used during training or validation. The target prices are analytical Black-Scholes prices, so the reported errors measure how accurately the network approximates the known pricing function.

The Monte Carlo benchmark is evaluated against the same analytical Black-Scholes reference, but only on a deterministic subset of 512 test options. This restriction keeps the numerical benchmark computationally bounded while still providing a useful comparison with a classical simulation-based method.

The main metrics are mean absolute error, root mean squared error, coefficient of determination, and MAPE computed only on prices greater than 1. The MAPE restriction avoids unstable percentage errors for options whose theoretical price is close to zero.

5.4 Diagnostic Visualizations

The experiment produces the following diagnostic plots:

- training loss and validation loss;
- comparison between the Black-Scholes price and the neural network prediction;
- distribution of prediction errors;
- absolute error as a function of moneyness S_0/K ;
- absolute error as a function of maturity T ;
- absolute error as a function of volatility σ .

6 Results

6.1 Neural Network Approximation

The final neural network results on the held-out test set are reported in Table 6.1.

Metric	Value
MAE	0.06198
RMSE	0.08219
R^2	0.999987
MAPE, prices > 1	0.6746%

Table 6.1: Final neural network metrics against analytical Black-Scholes prices.

The results show that the feed-forward neural network approximates the Black-Scholes pricing function with very high accuracy on the sampled parameter domain. The coefficient of determination is close to one, while MAE and RMSE are small in absolute price units. This supports the main research question: in this controlled setting, a neural network can learn the Black-Scholes pricing map effectively and can be interpreted as a fast pricing surrogate after training.




Figure 6.1: Training loss and validation loss for the final experiment.

Figure 6.1 shows that the optimization process remains stable. The validation curve follows the training curve closely, suggesting that the model does not simply memorize the training set.




Figure 6.2: Analytical Black-Scholes prices against neural network predictions.

Figure 6.2 confirms that predicted prices are tightly aligned with analytical prices across the observed price range.

6.2 Error Diagnostics




Figure 6.3: Distribution of neural network prediction errors.

The error distribution in Figure 6.3 is concentrated near zero, which is consistent with the low MAE and RMSE values reported in Table 6.1.




Figure 6.4: Absolute neural network error against selected financial features.

Figure 6.4 is used to inspect whether the approximation error changes systematically across the parameter space. This is important because aggregate metrics alone may hide localized weaknesses, especially around different moneyness regimes, short maturities, or high-volatility regions.

6.3 Monte Carlo Benchmark

The Monte Carlo benchmark is evaluated against the analytical Black-Scholes price on 512 deterministic test options. The results are shown in Table 6.2.

Metric	Value
MAE	0.08882
RMSE	0.15790
R^2	0.999951
MAPE, prices > 1	0.6483%

Table 6.2: Final Monte Carlo metrics against analytical Black-Scholes prices.

The Monte Carlo benchmark also achieves high accuracy, as expected when the number of simulated paths is large. However, it remains a simulation-based estimator and therefore has sampling error. The neural network, in contrast, produces deterministic forward-pass predictions once trained. This highlights the main computational role of the network: it does not replace the analytical Black-Scholes formula in this setting, but it demonstrates how a learned surrogate can approximate a pricing map efficiently.

6.4 Discussion

The final results indicate that the neural network performs slightly better than the Monte Carlo benchmark in terms of MAE and RMSE on the reported experiment, while Monte Carlo obtains a slightly lower MAPE on prices greater than 1. This comparison should be interpreted carefully because the two methods are evaluated differently: the neural network is evaluated on the full held-out test set, whereas Monte Carlo is evaluated on a deterministic subset due to computational cost.

The most important conclusion is qualitative rather than competitive. In a controlled Black-Scholes environment, a feed-forward network can learn the mapping from option and market parameters to theoretical prices with high accuracy. This validates the use of the Black-Scholes setting as a clean benchmark for studying neural-network pricing surrogates before moving to more complex models.

7 Conclusions and Next Steps

The final experiment confirms that a feed-forward neural network can approximate the Black-Scholes pricing function for European call options with high accuracy in a controlled synthetic environment. With 100,000 synthetic contracts, the network achieves an MAE of 0.06198, an RMSE of 0.08219, and an R^2 of 0.999987 on the held-out test set. These results answer the research question positively within the assumptions and parameter ranges considered in the project.

The comparison with Monte Carlo also clarifies the role of the neural network. Monte Carlo is a classical numerical estimator and remains useful as a benchmark, while the trained neural network acts as a deterministic pricing surrogate. The purpose of the network is therefore not to replace the analytical Black-Scholes formula, but to demonstrate that a supervised model can accurately learn a well-defined pricing map.

The main limitations are:

- the data are synthetic and generated under Black-Scholes assumptions;
- only European call options are considered;
- the target pricing function is known analytically;
- generalization outside the sampled parameter ranges is not guaranteed.

Natural next steps are:

- compare different neural network architectures and activation functions;
- study the error in specific regions of the parameter space, for instance deep in-the-money or out-of-the-money options;
- compare evaluation times across the analytical formula, Monte Carlo simulation, and the neural network;
- extend the framework to stochastic volatility models such as Heston;
- estimate Greeks using either automatic differentiation or finite differences.

These extensions are outside the scope of the current implementation, but they represent natural directions for generalizing the approach beyond the Black-Scholes benchmark.

Bibliography

- [1] Fischer Black and Myron Scholes. The pricing of options and corporate liabilities. *Journal of Political Economy*, 81(3):637–654, 1973.
- [2] George Cybenko. Approximation by superpositions of a sigmoidal function. *Mathematics of Control, Signals and Systems*, 2:303–314, 1989.
- [3] Kurt Hornik, Maxwell Stinchcombe, and Halbert White. Multilayer feedforward networks are universal approximators. *Neural Networks*, 2(5):359–366, 1989.
- [4] Steven E. Shreve. *Stochastic Calculus for Finance II: Continuous-Time Models*. Springer, 2004.